CS150 Final Project Report, Team 07

Sahar Mesri, Sagar Karandikar, `cs150-bw, cs150-bn`

December 11, 2013

Abstract

In this report we describe our hardware implementation of a special purpose pipeline for real-time edge-detection computation within a live VGA video stream. The hardware that accompanies our pipeline allows us to select between a variety of output streams, each of which originate from different stages in our two octave processing pipeline. As a result, the user is able to view the presence of edges in the input stream at progressively decreasing granularities. Additionally, the user is able to view the input stream with various degrees of gaussian filtering applied. Ultimately we attempt to recreate some of the functionality of the SIFT algorithm described in "A Parallel Hardware Architecture for Scale and Rotation Invariant Feature Detection" [Bonato, et. al].

1 Overview

1.1 Design - Full Pipeline Black-Box Block Diagram

See Figure 1.

1.2 Brief Description of Major Submodules

1.2.1 SRAM Arbiter

The SRAM arbiter is the main interface between SRAM and the rest of our system. Essentially, its task is to mediate reads and writes between other modules that are competing for access to our single port SRAM. The arbiter supports two write ports, two read ports, and services a single request (either a single read or a single write) per clock cycle.

1.2.2 Check4.v, Check4_4x.v

This module connects the entire Downsampler, Octave, and Upsampler pipeline together. Additionally, it provides the switching functionality that allows our design to display any of the gaussians or differences on-demand. See Figure 2 for a detailed connection diagram.

1.2.3 DownsamplerWrap.v, Downsampler4xWrap.v

This module attaches the Downsampler to its output fifo, which mediates between the Downsampler and the gaussian pipeline. This module spans both the first and second clock domains. A detailed block diagram can be found in Figure 3.
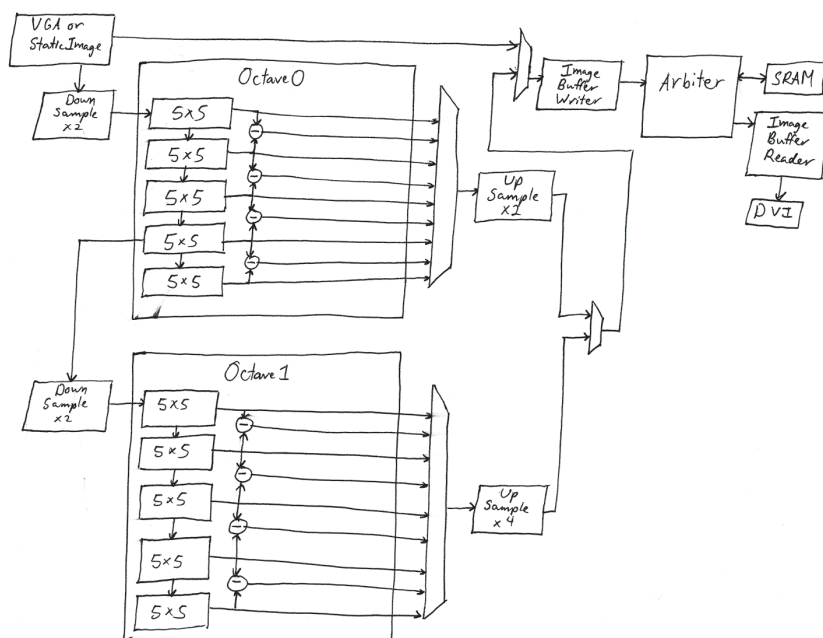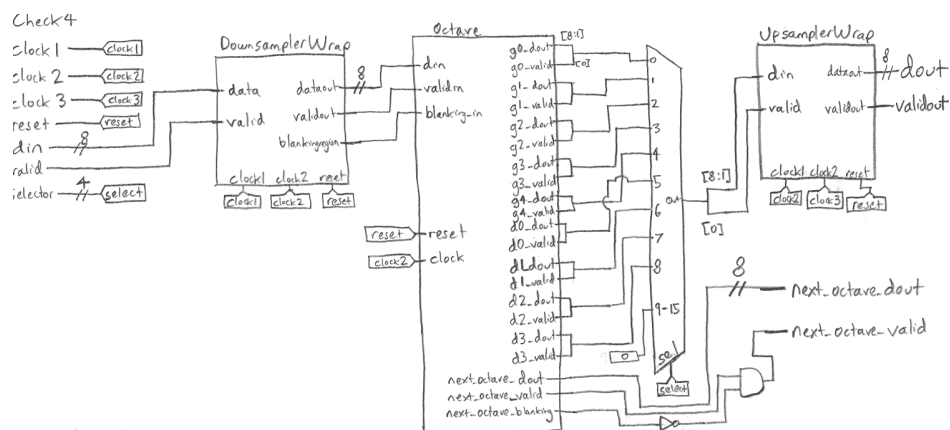
Figure 1: Full-Pipeline Blackbox Image


Figure 2: Check4.v Block Diagram

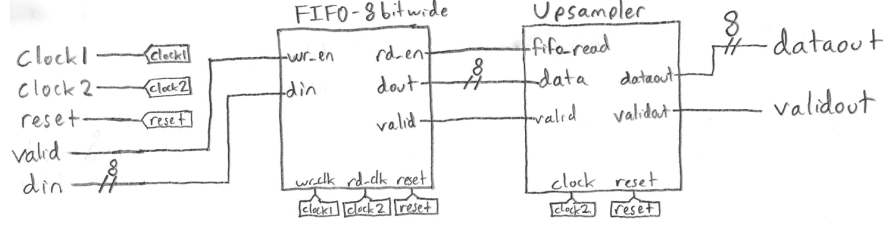Figure 3: DownsamplerWrap.v Block Diagram

Figure 4: UpsamplerWrap.v Block Diagram

1.2.4 UpsamplerWrap.v, Upsampler4xWrap.v

This module attaches the Upsampler to its input fifo, which mediates between the gaussian pipeline and the Upsampler. This module spans both the second and third clock domains. A detailed block diagram can be found in Figure 4.

1.2.5 Downsampler.v, Downsampler4x.v

The Downsampler module converts the original 800x600 image fed into our processing pipeline into a 400x300 image to align with our hardware constraints. The downsampler takes in the image data as a stream of pixels along with a valid signal that is only raised when the input pixel is valid data from the image. In addition to reducing image size, the downsampler is responsible for "generating" the blanking regions (sequences of black pixels) that are used to flush data through our gaussian pipeline. The downsampler ultimately outputs a scaled image as a stream of pixels along with a valid signal and a blankingregion signal that indicate whether the current output pixel is valid and/or a pixel in the blanking region respectively. Downsampler4x.v is identical to Downsampler.v, excluding some constants that are modified to allow this version to downsample from 400x300 images to 200x150 images. A black-box diagram of the Downsampler can be found in Figure 3.

1.2.6 Upsampler.v, Upsampler4x.v

The Upsampler module converts the processed 400x300 image coming out of the processing pipeline back into a 800x600 image for display over the DVI interface. The upsampler takes in the processed image data as a stream of pixels along with a valid signal that is asserted only for data that is both valid pixel data from the image and not part of the blanking regions. It performs basic upscaling by effectively taking each pixel from the processed input and expanding it to cover a 2x2 box of pixels in the final output. The upsampler outputs a pixel at a time in the output image along with a validout signal that is asserted whenever the data on the output is valid pixel data that we ultimately wish to display (blanking regions are non-existant at this point). Upsampler4x.v is identical to Upsampler.v, excluding some constants that are modified to allow this version to upsample from 200x150 images to 800x600 images. A black-box diagram of the Upsampler can be found in Figure 4.

1.2.7 octave.v

The Octave module contains 5 Gaussian filter blocks and computes the four Difference of Gaussian outputs. The module takes in a stream of pixels (one at a time), along with a valid signal and a blanking signal. It will output a stream of pixels (along with a valid signal) for each Gaussian filter and each Difference of Gaussian. There is no blanking output since

blanking pixels are considered invalid at the output. This module also has a special output just for the next octave, which contains the data, valid, and blanking signal from one of the Gaussian filters (the fourth one in this case). A black-box diagram of the Octave module can be found in Figure 2.

1.2.8   five_by_five_window.v

This module represents one Gaussian filter block, with five coefficients in the horizontal direction and five coefficients in the vertical direction. The input is a stream of pixels (one at a time), along with a valid signal and a blanking signal. The output is also a stream of pixels (after going through the Gaussian filter), along with a valid signal and a blanking signal. Pixels that are marked as blanking are also replaced with zero prior to the output. A black-box diagram of the five by five window can be found in Figure 5.

1.2.9   x_window.v

This module is used to compute the horizontal portion of the Gaussian filter. It will multiply 5 pixels in the same row by their corresponding coefficients. It will then take the average of those results to find the output value, whose location within the frame corresponds to the center of those five pixels. The input is a stream of pixels (one at a time), along with a valid signal and a blanking signal. The output is also a stream of pixels (after computing the horizontal Gaussian), along with a valid signal and a blanking signal. See Figure 6 for a black-box diagram.

1.2.10   five_row_array.v

This module is used to store five rows of pixels at a time, which becomes the input into the y_window module. This module takes in a stream of input pixels (one at a time), along with a valid signal and a blanking signal. It outputs five pixels at a time, one from each line, along with a valid signal and a blanking signal. These five pixels have the same horizontal position in the frame, and come from consecutive lines within the frame. The center of these five consecutive lines is the pixel at the center of the Gaussian filter block, so the blanking output corresponds to the blanking input for this particular pixel. See Figure 6 for a black-box diagram.

1.2.11   y_window.v

This module computes the vertical portion of the Gaussian filter. It multiplies 5 input pixels (from the same column) by their corresponding coefficients, and averages those results to find the output value. The location of the output value within the frame corresponds to the center of the input pixels. The input is a stream of pixels (five at a time), along with a valid signal and blanking signal. The output is also a stream of pixels (after computing the vertical Gaussian), along with a valid signal and a blanking signal. See Figure 6 for a black-box diagram.

2 System Description

2.1 Datapath

2.1.1 Connections to SRAM

All interfacing with the SRAM in our design is done solely through the SRAM Arbiter designed in Checkpoint 2. This Arbiter follows a predefined set of transition rules for determining order-of-service for the two write and two read ports that we will mention shortly. Although we would traditionally use a state diagram to describe this, for this module, a state-diagram is prohibitive since it simply consists of a complete graph over the 5 states (every possible transition exists). The five states that the SRAM arbiter can enter are the servicing states DOW0, DOW1, DOR0, and DOR1 and a PAUSE state in which no requests are serviced. Excluding the PAUSE state, state names can be decoded in the following manner: the last two characters stand for the port that is being serviced while the arbiter is in that state. If we order the servicing states as they are presented above, we can see the priority given to state transitions in our FSM. For example, if we are currently in state DOW1, we will service the R0 port if a read requests exists there, the R1 port if a read requests exists there and not at R0, the W0 port if neither R0 or R1 have a request and W0 does and finally the W1 port again if there a request exists and no other ports have a request. In any of our servicing states, if none of the ports have pending requests, we enter the PAUSE state, which follows the same priority as the W1 state would. In this way, we maintain fairness of access to the SRAM, allowing different parts of our pipeline to interact with the SRAM throughout execution.

2.1.2 Difference of Gaussian Filter Blocks

octave.v   This module is where the each Difference of Gaussian is computed for a single octave. It contains five Gaussian filter blocks, where the output of one block is fed into the input of the next block. The time between the first valid pixel of one block and the first valid pixel of the next block is determined by the time it takes for one pixel to travel through an entire Gaussian filter. To compute a Difference of Gaussian, the same pixel from each Gaussian filter block needs to be used, despite the delay between the two blocks. This means that the output of the first Gaussian must be stored in a shift register with a depth equal to that delay, while also being routed to the next Gaussian filter block. The output of the next Gaussian is then subtracted from the output of the shift register. The result is the Difference of Gaussian. In order to make the Difference of Gaussian more visible, the result is also multiplied by eight (implemented as a left shift by three).

The blanking signal also needs to be delayed, however the second of the two Gaussians already takes care that, so the Difference of Gaussian simply uses the blanking signal of the latter Gaussian filter block. Blanking signals are used by other modules within the octave, and indicate whether or not a particular pixel is part of the original image (not blanking), or was added in order to pad the edges of the image (blanking). Blanking pixels are not used by any module outside of the octave, so blanking pixels are simply considered invalid pixels at the output.

five_by_five_window.v   This module is one Gaussian filter block. It is split into three sub-modules: x_window, five_row_array, and y_window. This module serves four main purposes: 1) To connect x_window, five_row_array, and y_window to each other and to the inputs and outputs of the five_by_five_window module 2) To generate control signals for five_row_array
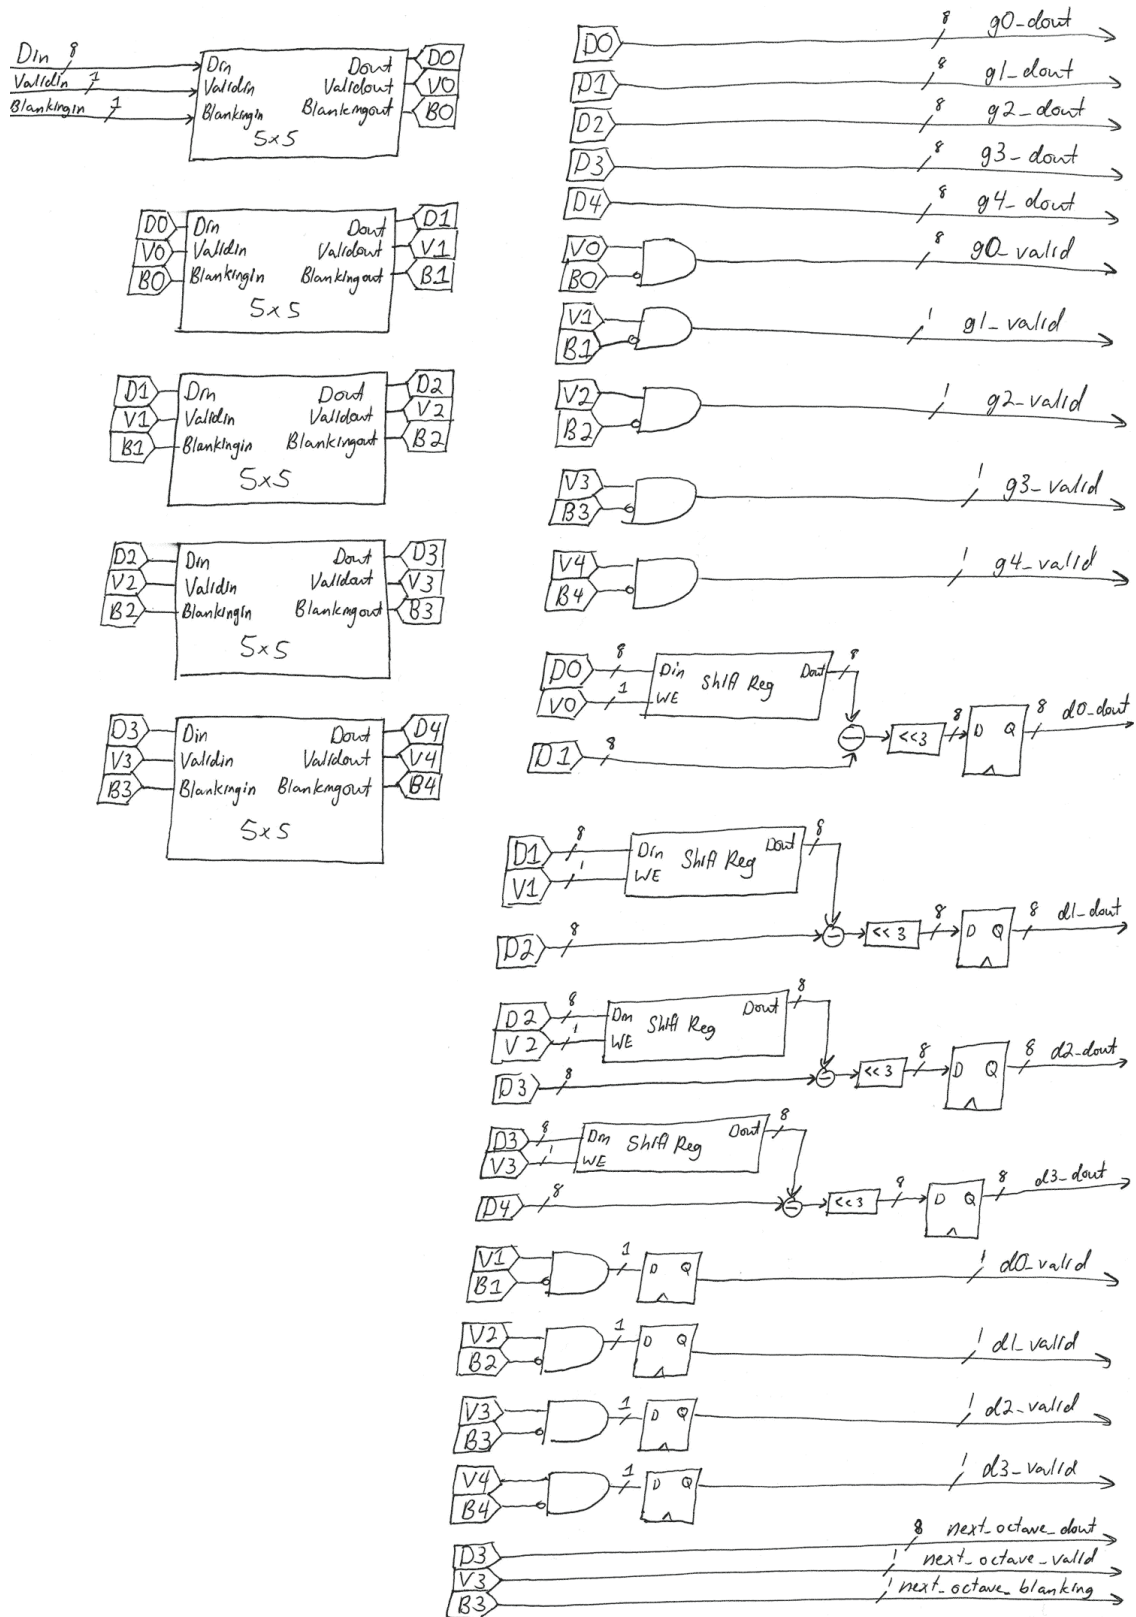
Figure 5: Full Octave Block Diagram

Figure 6: 5x5 Window Block Diagram

and y_window 3) To delay the blanking input and ensure a proper blanking output 4) To replace any blanking input pixels and output pixels with zero x_window, five_row_array, and y_window were designed so that they would connect directly to each other, in that particular order. The valid input into the module is fed straight into x_window, and the valid output of y_window is attached directly to the valid output of the five_by_five_window. The three sub-modules do not need the blanking input signal. This signal is instead put into a shift register, whose depth is equal to the delay it takes for one pixel to travel through x_window, five_row_array, and y_window. The output of this shift register is the blanking output signal.

The data into the five_by_five_window module first goes into a mux, which chooses between the input data and a zero contestant, depending on if the input is part of a blanking region or not. The output of that mux is then fed directly into x_window. In the same way, the output of y_window goes into a mux, which chooses between the output data or a zero constant, depending on if the output is blanking. The output of this mux is connected to the data output of the five_by_five_window module. This replacement of blanking pixels is done because those blanking pixels are used to pad the edges of the image, which is needed to keep the image size constant while feeding it through several Gaussian filters. In particular, the edges of the image need to be padded with zeros. However, running those zero pixels through an averaging filter will make them non-zero, which is why they need to be replaced by zeros between every filter block. The control signals generated by this module for five_row_array and y_window will be discussed in a later section.

x_window.v This module computes the horizontal portion of the Gaussian filter. The basic idea is that the module will store the 4 previous input pixels and the current pixel. It will also multiply each pixel by its coefficient, add the result of those operations, and scale the result back down to an 8-bit value for the next stage.

The module expects the value of each pixel in the frame as input, one line at a time (top to bottom), starting from the first (leftmost) pixel in the line and ending at the last
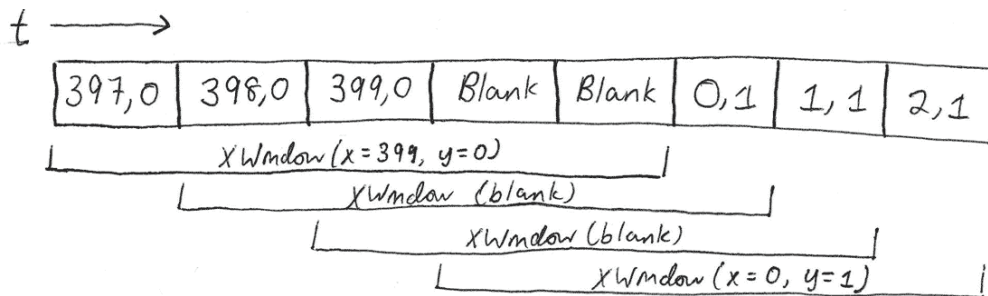
Figure 7: X Window Block Diagram

Figure 8: Example of Serial Input

(rightmost) pixel in the line. This module does not take into account padding. If the image is inputted without extra padding pixels, then the results at the edges of each line may not be accurate. To ensure proper results, there needs to be at least two blanking pixels at the end of this line (assuming the horizontal Gaussian window takes in five pixels). The two pixels can act as padding for both the previous line and the next line. An illustration of this concept is shown below:

Since all registers are set to zero upon reset, including the registers that store previous pixel values, the first two pixels of the frame are already padded.

In order to ensure maximum clock frequency, there is a register between each consecutive mathematical operation, such as a multiplication and addition. Because of this, computing the sum of all five terms takes several cycles, and adds some latency to this portion of the pipeline. However, once the pipeline is filled, there will be one valid output during every cycle with a valid input (invalid inputs will stall the pipeline). The time it takes to generate an output for a particular input (assuming the pipeline is not stalled after this point) is 7 cycles. This is also the time it takes to fill the pipeline. Therefore, there is a 7 cycle delay between when the valid input signal becomes high and when the valid output signal becomes high (again, assuming there are no stalls).
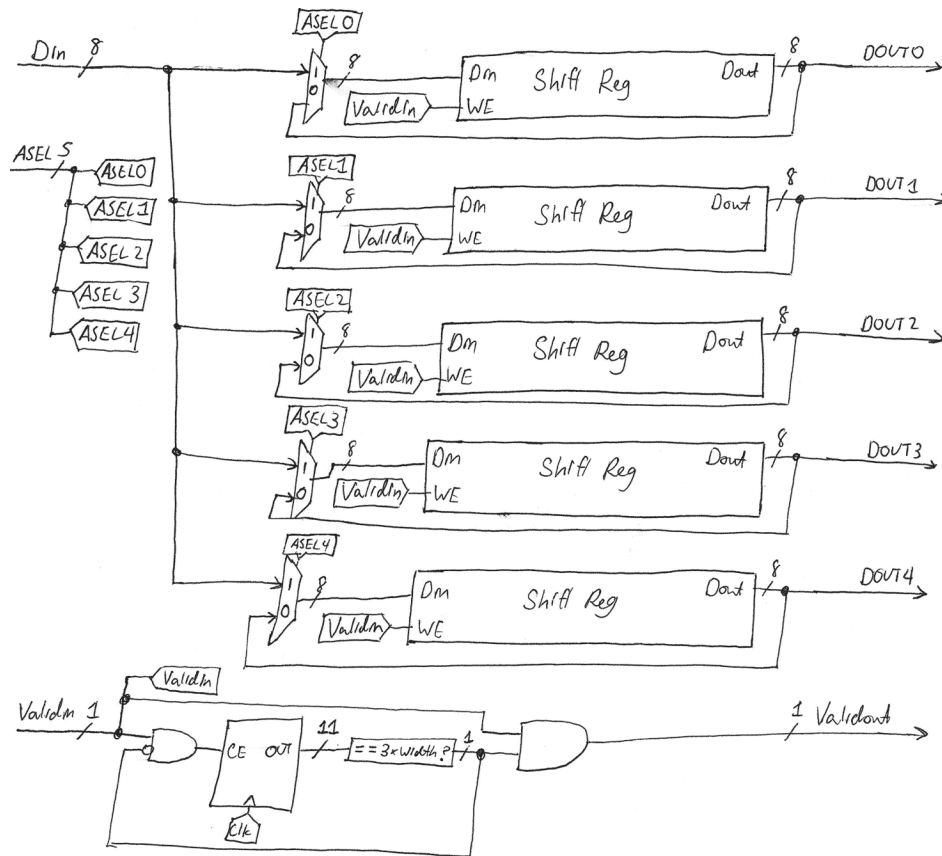
Figure 9: Five by Five Array Block Diagram

This module also takes advantage of the symmetry in the Gaussian coefficients, and computes the product of a pixel and its coefficient only once. For example, once the product of some pixel and the coefficient h1 is computed for the first time, that value will be needed again two cycles later. Likewise, the product of some pixel and the coefficient h0 will be needed 4 cycles later. The module will store and shift this value down the pipeline where it will be used for later operations which require this result. This reduces the number of multipliers in the module, which reduces the overall number of DSP SLICEs required for this particular module.

The coefficients are scaled in such a way that the result of adding each term (the product of a pixel and its coefficient) is no larger than a 16-bit unsigned number. In order to scale this back down to an 8-bit number, we simply keep the upper 8 bits and drop the lower 8. This is the same as dividing the result by 256, or shifting the result to the right by 8.

five_row_array.v　This module stores the results of the x_window module. Each output of x_window is the result of running each pixel through the horizontal portion of the Gaussian filter. This module will store 5 lines worth of pixels, which is needed to compute the vertical portion of the Gaussian filter. The output is five pixel values, all from the same position within a line (column), from five consecutive rows. This is accomplished by creating a 5 shift registers, one for each line, with a depth equal to the number of pixels in a line (including blanking). The module is constantly shifting values in and out of each shift register, as long as there is a valid input. One of the five shift registers gets new values from the input, and

every other shift register has its output fed back into the input. This is done to ensure that a pixel of a particular position within a line (such as x=250) is in the same position within all 5 shift registers. The array containing the oldest values gets replaced by newer values. Once that array is full, the module will begin filling/replacing the contents of the next oldest array. An illustration of this concept is shown below:

The input signal asel is used to choose which array to shift the new value into. This is a five-bit, one hot signal. Each bit corresponds to a particular shift register, and if that bit is 1 then the shift register will shift in the new data (instead of shifting in the old data). As long as each new value is shifted into the proper register in the proper order, then the outputs will also be in the proper order. See the example above for a demonstration on how this is accomplished.

Only the first three rows need to be filled before there is a valid output. This is because the shift registers are set to 0 upon reset, so two of the rows can act as blanking rows. This means that the delay between a valid input and a valid output (assuming no stalls/invalid inputs) is 3*width, where width is the number of pixels in a row, including blanking. This is also the time it takes to fill this portion of the pipeline, and the time it takes to generate an output for one particular input.

Note that the first shift register does not always correspond to the first line within the vertical Gaussian window (nor does the second shift register correspond to the second line, and so on). This is because the current line being filled with new data is always rotating (according to the input asel). This means that the y_window module needs to keep track of which line is which, and select the appropriate coefficients for each line. How and where these control signals are generated will be discussed in a later section.

y_window.v This module computes the vertical portion of the Gaussian filter. It will receive five input pixels from the five_row_array. Each pixel is from the same position in a line (such as x=250), and comes from consecutive lines. It is expected that the pixels will start from the leftmost position within a line, and end at the rightmost position. It is also expected that the groups of lines will be received in order from top to bottom (i.e. lines [1,2,3,4,5]followed by lines [2,3,4,5,6]). These pixels are then multiplied by their corresponding coefficients and the results are added together.

Figure 10: IO Diagram

This sum will then be scaled back down to an 8-bit number. To ensure proper results, at least two extra lines of zeros (blanking/padding), need to be added at the end of each frame. This is similar to the requirement of two extra pixels in each line for the x_window module.

In order to ensure maximum clock frequency, there is a register between each consecutive mathematical operation, such as a multiplication and addition. Because of this, computing the sum of all five terms takes several cycles, and adds some latency to this portion of the pipeline. However, once the pipeline is filled, there will be one valid output during every cycle with a valid input (invalid inputs will stall the pipeline). The time it takes to generate an output for a particular input (assuming the pipeline is not stalled after this point) is 5 cycles. This is also the time it takes to fill the pipeline. Therefore, there is a 5 cycle
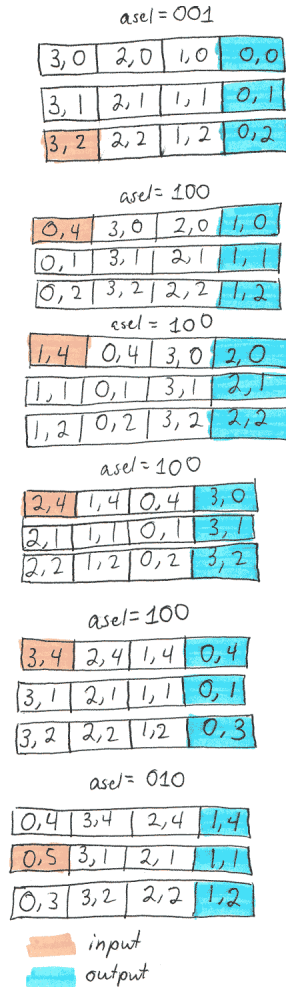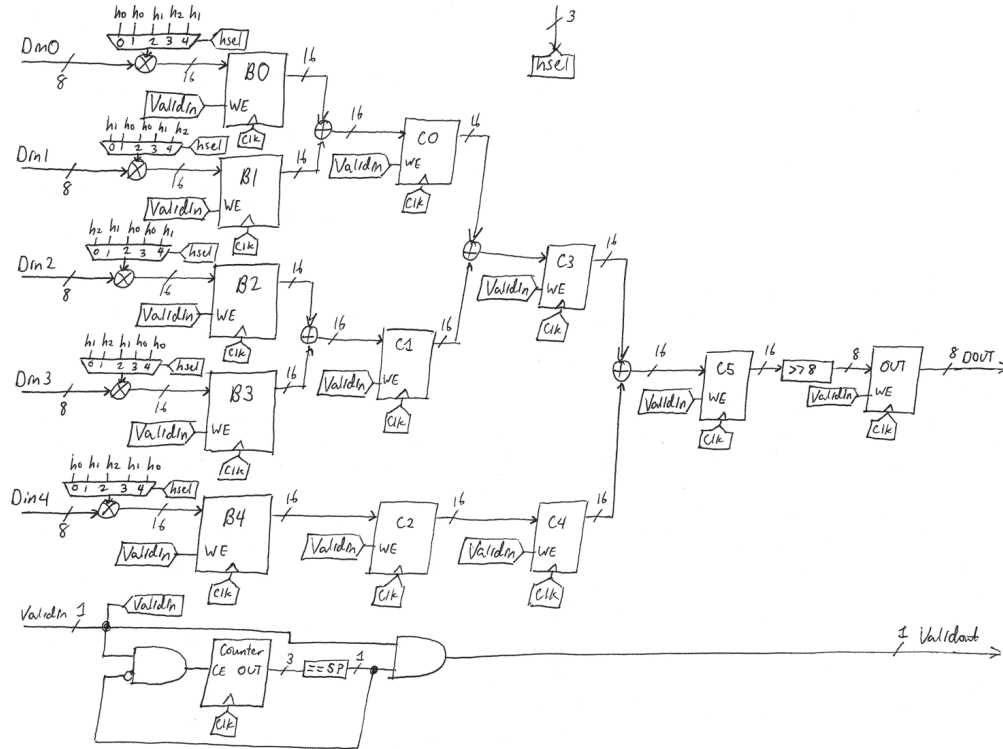
Figure 11: Y Window Block Diagram

delay between when the valid input signal becomes high and when the valid output signal becomes high (again, assuming there are no stalls). This is shorter than the 7 cycle delay of the x_window module because this module receives all 5 pixels at once, rather than having to store each pixel one at a time until there is enough for a valid output.

As we pointed out in the previous section, the outputs of the five_row_array do not always correspond to the same line within the vertical Gaussian window. This means that the inputs to y_window, which are connected directly to the outputs of five_row_array, also do not correspond to the same line. This also means that the coefficients (that we multiply each input by), cannot be constant. The signal hsel, which is a 3-bit number ranging between zero and four, is used to choose the coefficient for each line. In terms of hardware, this means that there is a 5-mux (for each line) which chooses between the five coefficients, and the result is then fed into the multiplier (for each line). The same signal hsel is used for all five multiplexers. What is different is the order in which the coefficients are attached to each multiplexer. This allows one signal, hsel, to choose between the 5 different ways in which the outputs from five_row_array can be arranged. As long as hsel and asel are in synch, and change at the same time, the proper coefficients will be used. Details on how hsel and asel are generated will be discussed in a later section.

The coefficients are scaled in such a way that the result of adding each term (the product of a pixel and its coefficient) is no larger than a 16-bit unsigned number. In order to scale this back down to an 8-bit number, we simply keep the upper 8 bits and drop the lower 8. This is the same as dividing the result by 256, or shifting the result to the right by 8.

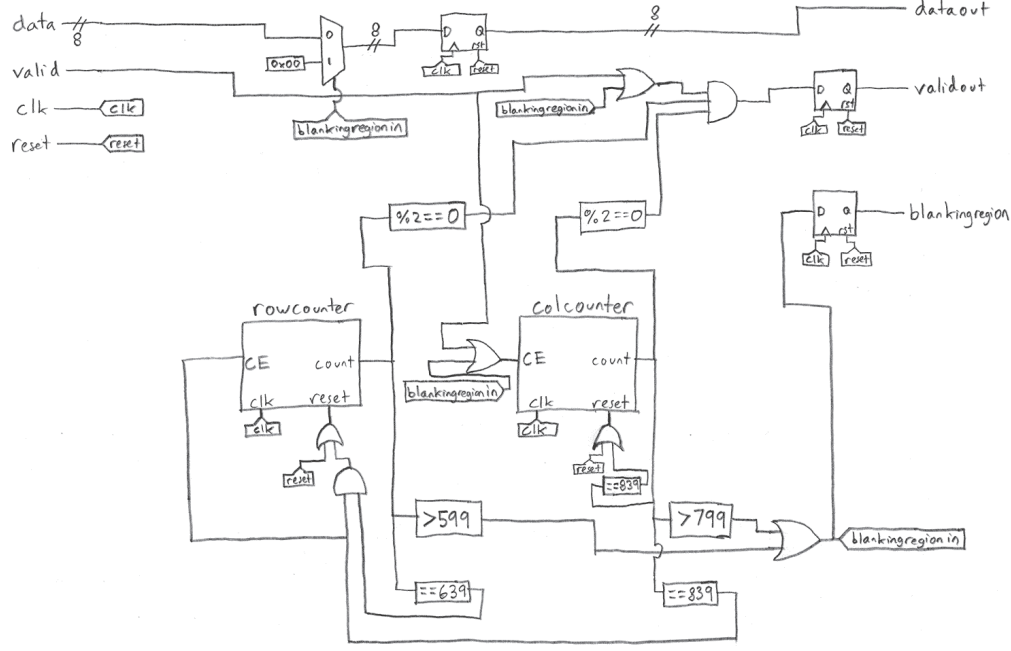## 2.1.3 Connections to ImageBufferWriter

Figure 12: Downsampler Block Diagram

ImageBufferWriter only requires a data input and a valid signal, and this comes from the upsampler. In this case, we have two upsamplers, one from each octave. Each octave also has nine outputs: five Gaussian outputs

Table 1: NOTE: Any missing binary numbers (switch combinations) are considered illegal and will cause the swap controller to stall.

Binary	00000	00001	00010	00011	00100	00101	00110	00111	01000
Octave	0	0	0	0	0	0	0	0	0
Image	gauss0	gauss1	gauss2	gauss3	gauss4	diff0	diff1	diff2	diff3
Binary	10000	10001	10010	10011	10100	10101	10110	10111	11000
Octave	1	1	1	1	1	1	1	1	1
Image	gauss0	gauss1	gauss2	gauss3	gauss4	diff0	diff1	diff2	diff3

puts and four Difference of Gaussian outputs. This means that we need to choose which one of those outputs goes into the upsampler, and then choose which upsampler to connect to ImageBufferWriter. The selector input for the Check4 and Check4_4x modules is used to choose between the five Gaussian outputs and the four Difference of Gaussian outputs. The value of selector is determined by GPIO_DIP[3:6]. Inside of FPGA_TOP_ML505, we use GPIO_DIP[7] to choose which upsampler to connect to ImageBufferWriter. If we consider GPIO_DIP[3:7] (a particular combination of switches) to be a 5-bit binary number, then the following table summarizes which image each number corresponds to:

Switching the input to ImageBufferWriter at the wrong time could result in a shifted image or a stalled swap controller. In order to mitigate these issues, the GPIO_DIP switches are not connected directly to the selectors/multiplexors. Instead, these values are stored in a register and that register is updated at a time when it is safe to switch. The details of this implementation and the timing will be discussed in a later section.

2.1.4 Downsampler

This module achieves two main data pre-processing goals for our system. Firstly, it scales the input image from 800x600 pixels per frame to 400x300 pixels per frame. Additionally, it provides "stable" blanking region generation in order to flush valid data from our pipeline
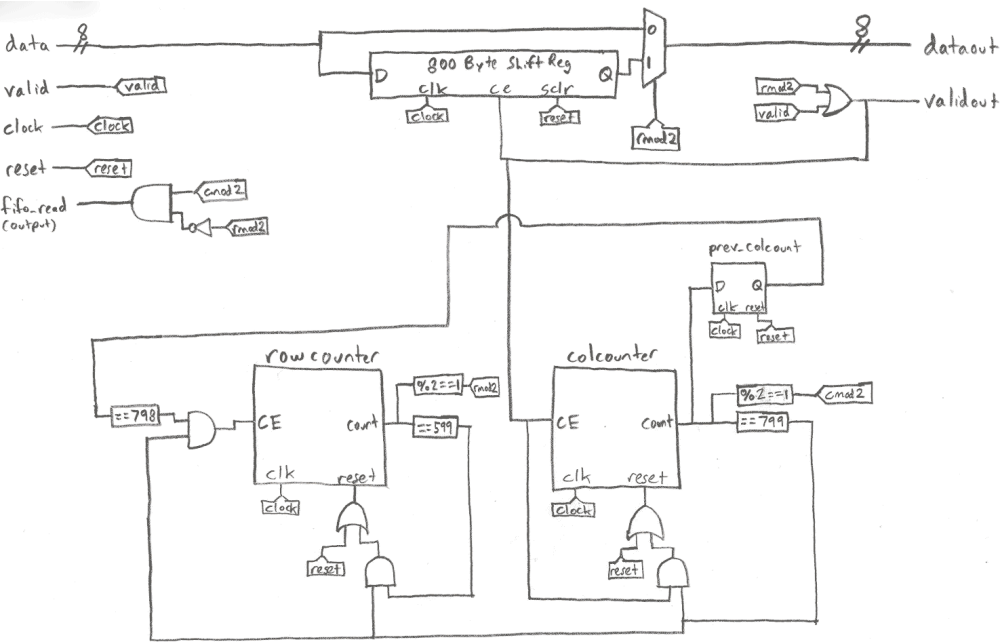
Figure 13: Upsampler Block Diagram

at the end of a line and the end of a frame. In doing so, it allows our pipeline to accept any input image that does not produce a valid output for at least 40 cycles at the end of a line and 40 lines at the end of a frame. These are reasonable values since they are far lower than the blanking regions generated by the VGA interface. The main feature in the datapath of the downsampler is a set of 3 registers, each tied to one of dataout (the 8 bit pixel output), validout (a 1-bit signal that is asserted if the output data is either part of the output image or part of a blanking region), and blankingregion (a 1 bit-signal that is asserted when the dataout value is a blanking region). The values fed into these registers are computed using the output of a set of two counters, which we will discuss in detail in the control section. Essentially, the validout signal is asserted for 420 cycles for every other row of the original image, plus 20 additional rows for blanking. The blankingout signal is similarly asserted for 20 cycles at the end of a row of the image and for 20 rows at the end of a frame of the image. This produces a 420x320 input image for the rest of the pipeline to process.

2.1.5 Upsampler

This module takes our processed 400x300 images and scales them up by 2x in each direction to produce an 800x600 image. The upsampler duplicates pixels using two methods. First, to duplicate in the horizontal direction, we simply instruct the fifo from which we are reading data to maintain each valid output on its data output line for exactly two cycles. Because we are utilizing the first-word-fall-through functionality of the FIFOs, we do not need to explicitly assert read enable in order to "get" values. Rather, the next value that the fifo is able to present is simply held at its output until we request the "next" value. Thus, upon seeing the assertion of the valid signal of the upsampler fifo, we simply wait one clock cycle before asserting our read enable signal to "clear out" the old value. Thus, as far as our output is concerned, we receive two copies of each valid pixel in the horizontal direction. Additionally, to produce the 2x duplication in the vertical direction, we take the output from our horizontal duplication "stage" and feed it into a shift register that is 800 bytes

wide. Thus, once we have processed an entire valid row of 400 input pixels (which outputs 800 valid pixels and additionally fills the shift register with 800 valid pixels), we switch our module's output to the output of the shift register and run for 800 cycles. This produces the duplicate row that we need for vertical scaling. This entire mechanism is controlled by a set of counters, which will be discussed in the control section.

2.2 Control

2.2.1 Dealing with issues - blank rows, addressing issues

Our pipeline propagates a valid signal throughout. This automatically invalidates any bad data that would otherwise flow into our pipeline (for example blank rows and bad data caused by stalls). Additionally, we do not face any addressing issues; as long as we feed exactly 480,000 pixels per frame, the ImageBufferWriter takes care of address generation automatically.

2.2.2 General controller design, by module

octave.v This module does not need any control signals or generate any control signals. However, it does depend on an input signal validin, and outputs a valid signal for each Gaussian and Difference of Gaussian. If there is no valid input, then the entire pipeline is stalled, and there are no valid outputs. When there are valid inputs, and the pipeline is still filling up, some Gaussian filter blocks will have a valid output and some will not. This is because the output of one Gaussian filter is fed into the input of the next Gaussian. This also means that the same pixel (for example, the first pixel), will be outputted at different times for different Gaussians. Once the entire pipeline has been filled, all outputs will be valid as long as there is a valid input as well. However, there is still a delay between the time when a particular pixel (i.e. a pixel at one specific location within the frame) shows up on the output of each Gaussian filter block. This delay is just the delay through all of the components of the Gaussian filter block (x_window, five_row_array, and y_window). The total delay is then (3*width)+12, where width is the number of pixels in one line, including blanking. Suppose width = 420, and the octave receives a constant stream of valid outputs starting at cycle 0. Then:

Note that the output for each Difference of Gaussian is delayed by one extra cycle. This is because there is a register after the subtraction and right shift operations to pipeline the output. Any stall in the pipeline (i.e. an invalid input into the octave for n number of cycles) will cause an additional delay.

Signal	Cycle Number
g0_valid	1272
g1_valid	2544
d0_valid	2545
g2_valid	3816
d1_valid	3817
g3_valid	5088
d2_valid	5089
g4_valid	6360
d3_valid	6361

Table 2: stuff

five_by_five_window.v This module does not need any control signals. However, it does generate the control signals asel and hsel for five_row_array and y_window, respectively. These signals are only supposed to change when the x_window module outputs enough pixels for one complete line. This was done by making a counter, which goes from 0 to width-1 (where width is the number of pixels in one line, including blanking), and increments when there is a valid output from x_window. The state machines which generate asel and hsel will change states whenever the counter is equal to width-1. Both state machines

are Moore machines which simply rotate between five different states. These states correspond to the five different ways in which the outputs of five_row_array can be arranged. The state diagrams for these machines are below:

This module also takes in a valid signal and outputs a valid signal. However, this modules does not do anything to these signals; it simply feeds the input into x_window, and takes the output directly from y_window. The behavior of this module's sub-modules (x_window, five_row_array, and y_window) does mean that this entire module will stall if there is no valid input.

Figure 14: ASEL and HSEL FSM Diagrams

x_window.v   This module does not need any control signals nor does it generate any control signals. It does have a simple counter to determine if the pipeline has filled up enough so that the output is valid. This counter will increment as long as there is a valid input, and the count is less than seven. Once the count has reached seven, the counter will stop incrementing unless the module is reset. The output of the module is invalid if the counter is less than seven, or the input is not valid. This also means that an invalid input will stall this module. A timing diagram demonstrating the relationship between validin, the counter, and validout is below:

five_row_array.v   This module does need one control signal, asel. What this signal is used for was discussed in the previous section on this module's datapath, and how this signal is generated was discussed in the previous section on five_by_five_window control signals. This module also does not generate any control signals. It does have a simple counter to determine if the pipeline has filled up enough so that the output is valid. This counter will increment as long

Figure 15: Timing Diagram

as there is a valid input, and the count is less than 3*width, where width is the number of pixels in a row (including blanking). Once the count has reached 3*width, the counter will stop incrementing unless the module is reset. The output of the module is invalid if the counter is less than 3*width, or the input is not valid. This also means that an invalid input will stall this module.
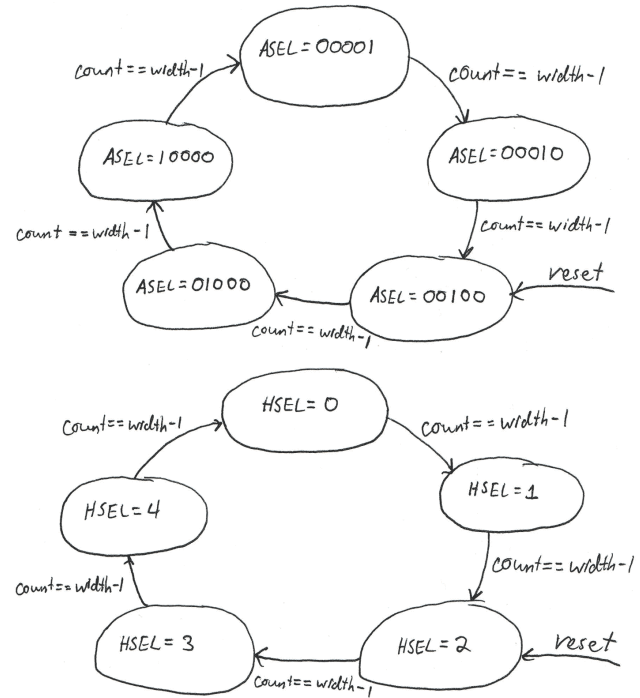
Signal	Boolean Expression
rmod4	rowcounter % 4 != 0
cmod4	colcounter % 4 == 3
fifo_read	cmod4 AND (NOT rmod4)
validout	rmod4 AND valid
rowcounter reset	colcounter == 799 AND rowcounter == 599
colcounter reset	validout AND colcounter == 799
rowcounter CE	prev_colcount == 798 AND colcount == 799
colcounter CE	validout

Figure 17: Upsampler (2x) Control / Intermediate Signal Summary

y_window.v This module does need one control signal, hsel. What this signal is used for was discussed in the previous section on this modules datapath, and how this signal is generated was discussed in the previous section on five_by_five_window control signals. This module also does not generate any control signals. It does have a simple counter to determine if the pipeline has filled up enough so that the output is valid. This counter will increment as long as there is a valid input, and the count is less than 5, where width is the number of pixels in a row (including blanking). Once the count has reached 5, the counter will stop incrementing unless the module is reset. The output of the module is invalid if the counter is less than 5, or the input is not valid. This also means that an invalid input will stall this module. A timing diagram demonstrating the relationship between validin, the counter, and validout is below:

Downsampler.v, Downsampler4x.v This module contains various control signals, all of which are driven by either a modulo or equality computation on a set of counters. These counters naturally represent our row number and column number in the frame currently being downsampled. These counters assume that the first valid we recieve is the first valid pixel in a frame. The counters then increment as we would expect when receiving a stream of pixels - the column


Figure 16: Timing Diagram

(horizontal) counter increments every cycle where valid is asserted, while the row (vertical) counter increments every cycle where the column counter has reached its maximum value (the width of the image minus one). Based on these values, the Downsampler generates what is effectively a control signal for the rest of the pipeline, its "validout" signal. For ease of reading, Table 3 presents the various computations done on the counters to genrate control / intermediate signals in the downsampler. The labels presented correspond with those in the detailed downsampler block diagram shown earlier.

Upsampler.v, Upsampler4x.v
This module contains various control signals, all of which are driven by either a modulo or equality computation on a set of counters. These counters naturally represent our row number and column number in the frame currently being upsampled. These counters assume that the first valid we receive is the first

Signal	Boolean Expression
blankingregionin	colcounter > 799 OR rowcounter > 599
rowcounter reset	colcounter == 839 AND rowcounter == 639
colcounter reset	colcounter == 839
rowcounterCE	colcounter == 839
colcounterCE	blankingregionin OR valid
validout	(colcounterCE) AND (rowcounter % 2 == 0) AND (colcounter % 2 == 0)

Table 3: Downsampler (2x) Control / Intermediate Signal Summary

valid pixel in the input frame. The counters then increment as we would expect when receiving a stream of pixels. In this case, the column (horizontal) counter increments twice for every valid input that the Upsampler sees, and 800 times after receiving a complete valid row. The row (vertical) counter always increments immediately after the cycle during which column counter transitions from (image width - 2) to (image width - 1). These counters are then used as indicated in Figure 17 to generate control signals for the rest of the module.

Image Switching   When switching inputs into the ImageBufferWriter, timing is critical. If the switch is made at the wrong time, the output might be shifted over, or the swap controller might get stalled. As mentioned in a previous section, GPIO_DIP[3:8] is the input to a register. The register is then connected to the mechanism that determines which image writes into ImageBufferWriter. All that is needed is a safe time to update that register. The best time to do this is right after the ImageBufferWriter has completed one frame. Once the ImageBufferWriter signals that it is finished, using bg_done, the swap controller will send bg_done_ack. Once that has happened, the ImageBufferReader, responsible for the DVI output, will get a swap signal from the swap controller. After the ImageBufferReader sends a swap_ack, the swap controller will issue a start signal (bg_start) to VGA or StaticImage, and wait for start_ack. StaticImage will return start_ack one cycle later, but VGA may have a delay (since it will only issue a start_ack at the end of a frame). The following timing diagram illustrates this process:

Note that start_ack may be delayed by one or more cycles if VGA is being used. If this happens, then bg_start will stay high for the same amount of cycles as the delay.

This gives us plenty of time to check if one of the switches has moved, reset the entire pipeline (including each octave, downsampler, and upsampler), and then switch the inputs into ImageBuffer-Writer. We need to reset the pipeline because some of the Gaussian filter blocks may still have some non-blanking pixels inside, which would be considered valid and would then be pushed into ImageBuffer-Writer. By resetting the pipeline, everything starts over, as if the next frame were

Figure 18: Timing Diagram

actually the first frame. We need all of this time because it takes more than one clock cycle to reset the FIFOs used in the upsampler and downsampler. According to the Xilinx Logi-CORE IP FIFO Generator v9.1 User Guide, the FIFOs we are using have an asynchronous reset, and it takes three clock cycles (write or read) after the asynchronous reset is detected on the rising edge read and write clock respectively. Since both the read and write ports of all of the FIFOs are on the same clock, that means it will take 3 clock cycles before the pipeline is actually ready to receive valid data. From the timing diagram above, we see that there will always be enough time to complete a reset. Below is an image, taken from the Xilinx LogiCORE IP FIFO Generator v9.1 User Guide, which demonstrates the behavior of a FIFO during a reset:

2.3 Design Decisions and Tradeoffs

2.3.1 Padding

As opposed to "resetting" the entire pipeline at the end of each line and frame, we chose to take advantage of the VGA blanking period to generate blanking regions (valid pixels with value equal to zero). These blanking regions serve two purposes: to provide zero padding for correct computation in the gaussian pipeline and to flush out any remaining valid pixels in the pipeline at the end of a line or frame. This method is advantageous in that it eliminates the latency for repeatedly refilling the processing pipeline during our gaussian computation.


Figure 19: Fifo Timing Diagram from Xilinx User Guide

2.3.2 Clock Rates

Our system utilizes two separate clock domains, although it is capable of supporting up to three. In our implementation the DVI interface, image buffer, SRAM, and the SRAM Arbiter all run on a clock of 50Mhz. The other domain, the bg_clock domain (running at a frequency based on the VGA source clock), drives all other elements of our processing pipeline.

Even though the design supports three separate clock domains due to initial concerns of a potential timing violation in the gaussian pipeline, the design in practice "ignores" the middle clock domain. This domain was initially intended to run at a slower frequency to allow more time for segments of the gaussian pipeline to complete. However, due to our extensive use of pipelining, the critical path was short enough such that a slower clock was not necessary. In our implementation, this middle domain is tied to the source (VGA) clock.

2.3.3 Moore vs. Mealy Machines

Because most of our state machines require knowledge of the position of a pixel in a given frame, we must be able to number each of the pixels inside of one such frame. Additionally, most of our state machines also need to know the position of a pixel within a line. These two constraints lead us to using a dual-counter based control methodology for each state machine (one row counter and one column counter). By nature of our counter FSMs (which apply mathematical operations on the current states to determine control signals), we use only Moore Machines in our design.

2.3.4 Issues Encountered During Implementation

Much of the issues that arose during the implementation phase were due to human error and timing issues as opposed to flaws in the original design. We relied mainly on authoring testbenches, confirming output in Modelsim, and physically testing our impelementation with the VGA display.

None of these issues required significant changes to our design. Most of our solutions involved minimal changes to the original implementations. Ultimately due to the random nature of these malfunctions, we made little use of ChipScope in our debugging. Most of the time were not able to capture sufficient amounts of data to determine the causes of issues using this technique.

3   Design Metrics

3.1   Timing

In a visual inspection of the Post-PAR Static Timing Report, we determined the critical path to be the following:

Location	Delay(ns)	Physical Resource Logical Resource(s)
SLICE_X43Y87.AQ	0.450	ch/down_to_five_valid ch/down/fifo2/U0/ xst_fifo_generator/ gconvfifo.rf/ grf.rf/gntv_or_sync_fifo. gl0.rd/gr1.rfwft/ user_valid
SLICE_X61Y65.D1	2.497	ch/down_to_five_valid
SLICE_X61Y65.D	0.094	ch/oct/.gauss0a/ my_x_window/ valid_count<2> ch/oct/.gauss0a/my_x_window/ validout1
SLICE_X79Y56.B1	1.858	ch/oct/.gauss0a/xvalidout
SLICE_X79Y56.B	0.094	ch/oct/.gauss0a/ my_five_row_array/ validout_cmp_eq0000 ch/oct/.gauss0a/my_y_window/ validout1
SLICE_X97Y61.A1	1.795	ch/oct/gauss0_validout
SLICE_X97Y61.A	0.094	ch/oct/.gauss1a/avalidout ch/oct/.gauss1a/ my_five_row_array/validout1
DSP48_X0Y55.CEP	6.868	ch/oct/.gauss1a/avalidout
DSP48_X0Y55.CLK	0.728	ch/oct/.gauss1a/my_y_window/ Mmult_B3_mult0001 ch/oct/.gauss1a/my_y_window/ Mmult_B3_mult0001
Total	14.478ns	(1.460ns logic, 13.018ns route) (10.1\% logic, 89.9\% route)

However, the minimum clock period reported to us by the tools was 17.060ns, which results in a maximum frequency of 58.617 MHz.

3.2 Utilization

- The number of LUTs used in our design was 14,442, for a utilization percentage of 20%.

- The number of DSP units used in our design was 64, for a utilization percentage of 100%.

- The number of BlockRAM/FIFOs used in our design was 13, for a utilization percentage of 8%.

From these specs, it appears that we can theoretically produce a design that contains 50 gaussian filter blocks (five times as large as our existing design). However, this is misleading since at this point we have exhausted all of the DSP resources available to us. Thus, any operation that would otherwise the use of a DSP unit (for example multiplication) now requires a large number of LUTs to implement. Thus the expected factor by which we can increase our design is actually far smaller.

3.3 Division of Labor

Upon receiving each checkpoint specification, we defined the modules and interfaces necessary to implement our design. From this point, we determined the expected amount of work for each module and attempted to balance the workload amongst ourselves.

Sahar - Main processing pipeline, including octaves and gaussians (Everything inside of the Octave).

Sagar - Support structures for the pipeline, including SRAM Arbiter, Down and Up-samplers, and overarching connections between the modules and FPGA_TOP. (Everything outside of the Octave).

Implementation time: 15 Hours in Design, 12 Hours in Implementation, 50 Hours in Debugging (including one all-nighter).

4 Conclusion

Our pipeline accepts an 800x600 image via a VGA input interface. This image is then down-sampled to a 400x300 image to fit within our memory size constraints. The downsampled image is then fed into a cascade of gaussian filter blocks. From this series of gaussian filter blocks, we derive outputs that are then subtracted to produce our difference of gaussians output. This DoG output is then fed into an upsampler to restore the original image size of 800x600 for display over the DVI interface. Additionally, the user can select between various gaussian and difference of gaussian outputs for display.

4.1 Ideas for Future Improvement

- Attempt a 7x7 Window.

- Implement the design without Downsampling and Upsampling.

- Switch roles (Implementing support structures instead of gaussian blocks and vice-versa).